

Semantic Kernel — Components (Concise Reference)

A short, visual-friendly document summarizing the Semantic Kernel components (from Microsoft Learn).

1. One-line overview

Semantic Kernel (SK) is an orchestration layer that lets you combine LLMs, memory (vector stores), plugins (functions), prompt templates and planners to build agentic applications.

2. Core components (quick bullets)

- **AI Service Connectors** — Abstracts Chat Completion / Text Generation / Embeddings / TTS / STT and other model services so SK can call different providers through a common interface.
 - **Vector Store (Memory) Connectors** — Abstraction for vector databases used for memory and retrieval (RAG). SK does not use them automatically; expose them as plugins or use via prompt templates.
 - **Plugins & Functions** — Named containers of functions. Sources:
 - Native code functions
 - OpenAPI-backed functions
 - Prompt templates (registered as functions)
 - Text-search (ITextSearch) adapters for RAGPlugins can be:
 1. Advertised to the chat model (so the model can request them)
 2. Called directly from templates or code
 - **Prompt Templates** — Combine context, instructions, user input and function output. Can be used as the starting point for a chat flow or registered as plugin functions. Templates can call other plugins and produce multi-step flows.
 - **Filters** — Hooks that run before/after events during the chat flow (before/after prompt rendering, before/after function invocation). Filters can be nested and are registered with the kernel.
-

3. How they fit together (flow diagram, simple)

```
User message
  ↓
Prompt Template (renders) → executes hardcoded plugin calls (if any)
  ↓
Kernel invokes Chat Completion AI (via AI Service Connector)
  ↓
Chat model may request Plugin functions (advertised) or return text
  ↓
Kernel (using Filters) runs pre/post hooks
  ↓
Optional: Vector Store used by a plugin for retrieval
  ↓
Final answer returned to user
```

4. Minimal C# snippet (registering services)

```
// PSEUDO-CODE (concise)
var kernel = Kernel.Builder.Build();

kernel.Config.AddOpenAIChatService("gpt-deploy", endpoint:
"<ENDPOINT>", apiKey: "<KEY>");

kernel.Memory.RegisterVectorStore("my-vs", new
AzureSearchVectorStore(...));

kernel.RegisterNativeFunction("utils", "Summarize", (input) =>
SummarizeText(input));

var result = await kernel.RunAsync("Hello, how can I help?");
```

5. Minimal Python snippet (registering services, plugins, vector store, filters)

```
import asyncio
import logging

from semantic_kernel import Kernel
from semantic_kernel.utils.logging import setup_logging
from semantic_kernel.functions import kernel_function
from semantic_kernel.connectors.ai.open_ai import AzureChatCompletion
from semantic_kernel.connectors.ai.function_choice_behavior import
FunctionChoiceBehavior
from semantic_kernel.connectors.ai.chat_completion_client_base import
ChatCompletionClientBase
from semantic_kernel.contents.chat_history import ChatHistory
from semantic_kernel.functions.kernel_arguments import KernelArguments
from
```

```
semantic_kernel.connectors.ai.open_ai.prompt_execution_settings.azure_
chat_prompt_execution_settings import (
    AzureChatPromptExecutionSettings,
)

async def main():
    # Initialize the kernel
    kernel = Kernel()

    # Add Azure OpenAI chat completion
    chat_completion = AzureChatCompletion(
        deployment_name="your_models_deployment_name",
        api_key="your_api_key",
        base_url="your_base_url",
    )
    kernel.add_service(chat_completion)

    # Set the logging level for semantic_kernel.kernel to DEBUG
    setup_logging()
    logging.getLogger("kernel").setLevel(logging.DEBUG)

    # Add a plugin (LightsPlugin should be defined below)
    kernel.add_plugin(
        LightsPlugin(),
        plugin_name="Lights",
    )

    # Enable planning
    execution_settings = AzureChatPromptExecutionSettings()
    execution_settings.function_choice_behavior =
FunctionChoiceBehavior.Auto()

    # Create a history of the conversation
    history = ChatHistory()

    # Chat loop
    while True:
        userInput = input("User > ")

        if userInput.lower() == "exit":
            break

        history.add_user_message(userInput)

        result = await chat_completion.get_chat_message_content(
            chat_history=history,
            settings=execution_settings,
            kernel=kernel,
        )

        print("Assistant > " + str(result))
```

```

        history.add_message(result)

# Run the main function
if __name__ == "__main__":
    asyncio.run(main())
```csharp
// PSEUDO-CODE (concise)
var kernel = Kernel.Builder.Build();

// Register AI service (Chat completion / text generation)
kernel.Config.AddOpenAIChatService("gpt-deploy", endpoint:
"<ENDPOINT>", apiKey: "<KEY>");

// Register a vector store connector (memory)
kernel.Memory.RegisterVectorStore("my-vs", new
AzureSearchVectorStore(...));

// Register a native function plugin
kernel.RegisterNativeFunction("utils", "Summarize", (input) =>
SummarizeText(input));

// Run a prompt template
var result = await kernel.RunAsync("Hello, how can I help?");

```

---

## 6. Practical notes / tips

- **Plugins are powerful:** expose data sources and business logic as plugin functions rather than embedding everything in prompts.
  - **Prompt templates as functions** enable language-defined behaviour (describe the function in natural language and register it).
  - **Use Filters** for logging, telemetry, auth, or modifying context before/after calls.
  - **Vector stores are explicit:** SK doesn't automatically search them; wire a retrieval plugin for RAG.
- 

## 7. Quick use-cases

- Knowledge-base assistant (KB search plugin + SK prompt templates)
  - Multi-step enterprise workflows (planner + plugins)
  - Personalization and session memory (vector store + memory connectors)
  - Tool-enabled agents (calendar, DB, search via OpenAPI plugins)
-